

```
In [1]: import pandas as pd
import numpy as np
```

```
In [2]: df_train=pd.read_csv("train1.csv")
df_test=pd.read_csv('test1.csv')
```

```
In [3]: df_train.head()
```

```
Out[3]:
```

	ID	y	X0	X1	X2	X3	X4	X5	X6	X8	...	X375	X376	X377	X378	X379	X380	X382	X383	X384
0	0	130.81	k	v	at	a	d	u	j	o	...	0	0	1	0	0	0	0	0	0
1	6	88.53	k	t	av	e	d	y	l	o	...	1	0	0	0	0	0	0	0	0
2	7	76.26	az	w	n	c	d	x	j	x	...	0	0	0	0	0	0	1	0	0
3	9	80.62	az	t	n	f	d	x	l	e	...	0	0	0	0	0	0	0	0	0
4	13	78.02	az	v	n	f	d	h	d	n	...	0	0	0	0	0	0	0	0	0

5 rows × 378 columns

```
In [4]: df_test.head()
```

```
Out[4]:
```

	ID	X0	X1	X2	X3	X4	X5	X6	X8	X10	...	X375	X376	X377	X378	X379	X380	X382	X383	X384
0	1	az	v	n	f	d	t	a	w	0	...	0	0	0	1	0	0	0	0	0
1	2	t	b	ai	a	d	b	g	y	0	...	0	0	1	0	0	0	0	0	0
2	3	az	v	as	f	d	a	j	j	0	...	0	0	0	1	0	0	0	0	0
3	4	az	l	n	f	d	z	l	n	0	...	0	0	0	1	0	0	0	0	0
4	5	w	s	as	c	d	y	i	m	0	...	1	0	0	0	0	0	0	0	0

5 rows × 377 columns

```
In [5]: df_train.shape
```

```
Out[5]: (4209, 378)
```

```
In [6]: df_test.shape
```

```
Out[6]: (4209, 377)
```

```
In [7]: y_train=df_train['y'].values
```

```
In [8]: y_train
```

```
Out[8]: array([130.81,  88.53,  76.26, ..., 109.22,  87.48, 110.85])
```

**TASK1-If for any column(s), the variance is equal to zero, then you need to remove those variable(s).**

```
In [9]: df_train.var()
```

```
C:\Users\Hp\AppData\Local\Temp\ipykernel_2640\3907004805.py:1: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.
```

```
df_train.var()
```

```
Out[9]: ID          5.941936e+06
        y           1.607667e+02
        X10         1.313092e-02
        X11         0.000000e+00
        X12         6.945713e-02
        ...
        X380        8.014579e-03
        X382        7.546747e-03
        X383        1.660732e-03
        X384        4.750593e-04
        X385        1.423823e-03
        Length: 370, dtype: float64
```

```
In [10]: df_train.var () ==0
```

```
C:\Users\Hp\AppData\Local\Temp\ipykernel_2640\2195562811.py:1: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.
```

```
df_train.var () ==0
```

```
Out[10]: ID          False
        y           False
        X10         False
        X11          True
        X12         False
        ...
        X380        False
        X382        False
        X383        False
        X384        False
        X385        False
        Length: 370, dtype: bool
```

```
In [11]: (df_train.var () ==0).values
```

```
C:\Users\Hp\AppData\Local\Temp\ipykernel_2640\4044553682.py:1: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.
```

```
(df_train.var () ==0).values
```

```
Out[11]: array([False, False, False, True, False, False, False, False, False,
False, False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False, False,
False, False, True, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, True, False,
False, False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False, False,
False, False, False, True, False, True, False, False, False, False,
False, False, False, False, False, False, False, False, False, False,
False, False, True, False, False, False, False, False, False, False,
False, False, False, False, False, False, True, True, False, False,
True, False, False, False, True, False, False, False, False, False,
False, False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False, False,
True, False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False, True,
False, False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False, False,
False, False, False, False, False, False, False, False, False, False,
False])
```

```
In [12]: variance_with_zero=df_train.var()[df_train.var()==0].index.values
variance_with_zero
```

C:\Users\Hp\AppData\Local\Temp\ipykernel\_2640\3026110461.py:1: FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric\_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.

```
variance_with_zero=df_train.var()[df_train.var()==0].index.values
Out[12]: array(['X11', 'X93', 'X107', 'X233', 'X235', 'X268', 'X289', 'X290',
'X293', 'X297', 'X330', 'X347'], dtype=object)
```

```
In [13]: df_train =df_train.drop(variance_with_zero, axis=1)
```

```
In [14]: df_train.shape
```

```
Out[14]: (4209, 366)
```

```
In [15]: df_train=df_train.drop('ID',axis=1)
```

```
In [16]: df_train.shape
```



# TASK3-Apply label encoder

```
In [20]: train_object_datasets=df_train.select_dtypes(include='object')
train_object_datasets
```

```
Out[20]:
```

	X0	X1	X2	X3	X4	X5	X6	X8
0	k	v	at	a	d	u	j	o
1	k	t	av	e	d	y	l	o
2	az	w	n	c	d	x	j	x
3	az	t	n	f	d	x	l	e
4	az	v	n	f	d	h	d	n
...	...	...	...	...	...	...	...	...
4204	ak	s	as	c	d	aa	d	q
4205	j	o	t	d	d	aa	h	h
4206	ak	v	r	a	d	aa	g	e
4207	al	r	e	f	d	aa	l	u
4208	z	r	ae	c	d	aa	g	w

4209 rows × 8 columns

```
In [21]: from sklearn import preprocessing
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
```

```
In [22]: df_train['X0'].unique()
```

```
Out[22]: array(['k', 'az', 't', 'al', 'o', 'w', 'j', 'h', 's', 'n', 'ay', 'f', 'x',
'y', 'aj', 'ak', 'am', 'z', 'q', 'at', 'ap', 'v', 'af', 'a', 'e',
'ai', 'd', 'aq', 'c', 'aa', 'ba', 'as', 'i', 'r', 'b', 'ax', 'bc',
'u', 'ad', 'au', 'm', 'l', 'aw', 'ao', 'ac', 'g', 'ab'],
dtype=object)
```

```
In [23]: df_train['X0'] = le.fit_transform(df_train['X0'])
```

```
In [24]: df_train['X0'].unique()
```

```
Out[24]: array([32, 20, 40,  9, 36, 43, 31, 29, 39, 35, 19, 27, 44, 45,  7,  8, 10,
46, 37, 15, 12, 42,  5,  0, 26,  6, 25, 13, 24,  1, 22, 14, 30, 38,
21, 18, 23, 41,  4, 16, 34, 33, 17, 11,  3, 28,  2])
```

```
In [25]: df_train['X1'] = le.fit_transform(df_train['X1'])
df_train['X2'] = le.fit_transform(df_train['X2'])
df_train['X3'] = le.fit_transform(df_train['X3'])
df_train['X4'] = le.fit_transform(df_train['X4'])
df_train['X5'] = le.fit_transform(df_train['X5'])
df_train['X6'] = le.fit_transform(df_train['X6'])
df_train['X8'] = le.fit_transform(df_train['X8'])
```

```
In [26]: df_train.head()
```

	y	X0	X1	X2	X3	X4	X5	X6	X8	X10	...	X375	X376	X377	X378	X379	X380	X382	X383	X384
0	130.81	32	23	17	0	3	24	9	14	0	...	0	0	1	0	0	0	0	0	0
1	88.53	32	21	19	4	3	28	11	14	0	...	1	0	0	0	0	0	0	0	0
2	76.26	20	24	34	2	3	27	9	23	0	...	0	0	0	0	0	0	1	0	0
3	80.62	20	21	34	5	3	27	11	4	0	...	0	0	0	0	0	0	0	0	0
4	78.02	20	23	34	5	3	12	3	13	0	...	0	0	0	0	0	0	0	0	0

5 rows × 365 columns

## TASK4-Perform dimensionality reduction.

```
In [27]: from sklearn import model_selection
from sklearn.model_selection import train_test_split
```

```
In [28]: X = df_train.drop('y', axis=1)
y = df_train.y
x_train, x_test, y_train, y_test = train_test_split(X,y,random_state=42,test_size=0.3)
```

```
In [29]: x_train.shape,x_test.shape,y_train.shape,y_test.shape
```

```
Out[29]: ((2946, 364), (1263, 364), (2946,), (1263,))
```

```
In [30]: from sklearn.decomposition import PCA
```

```
In [31]: sklearn_pca=PCA(n_components=0.95)
```

```
In [32]: sklearn_pca.fit(x_train)
```

```
Out[32]: PCA(n_components=0.95)
```

```
In [33]: x_train=sklearn_pca.transform(x_train)
```

```
In [34]: x_train.shape
```

```
Out[34]: (2946, 6)
```

```
In [35]: sklearn_pca=PCA(n_components=0.95)
sklearn_pca.fit(x_test)
```

```
Out[35]: PCA(n_components=0.95)
```

```
In [36]: x_test=sklearn_pca.transform(x_test)
x_test.shape
```

```
Out[36]: (1263, 6)
```

## pca for test\_df

```
In [37]: df_test.head()
```

```
Out[37]:
```

	ID	X0	X1	X2	X3	X4	X5	X6	X8	X10	...	X375	X376	X377	X378	X379	X380	X382	X383	X384	:
0	1	az	v	n	f	d	t	a	w	0	...	0	0	0	1	0	0	0	0	0	
1	2	t	b	ai	a	d	b	g	y	0	...	0	0	1	0	0	0	0	0	0	
2	3	az	v	as	f	d	a	j	j	0	...	0	0	0	1	0	0	0	0	0	
3	4	az	l	n	f	d	z	l	n	0	...	0	0	0	1	0	0	0	0	0	
4	5	w	s	as	c	d	y	i	m	0	...	1	0	0	0	0	0	0	0	0	

5 rows × 377 columns

```
In [38]: test_object_datasets=df_test.select_dtypes(include='object')
test_object_datasets
```

```
Out[38]:
```

	X0	X1	X2	X3	X4	X5	X6	X8
0	az	v	n	f	d	t	a	w
1	t	b	ai	a	d	b	g	y
2	az	v	as	f	d	a	j	j
3	az	l	n	f	d	z	l	n
4	w	s	as	c	d	y	i	m
...	...	...	...	...	...	...	...	...
4204	aj	h	as	f	d	aa	j	e
4205	t	aa	ai	d	d	aa	j	y
4206	y	v	as	f	d	aa	d	w
4207	ak	v	as	a	d	aa	c	q
4208	t	aa	ai	c	d	aa	g	r

4209 rows × 8 columns

```
In [39]: df_test['X0'].unique()
```

```
Out[39]: array(['az', 't', 'w', 'y', 'x', 'f', 'ap', 'o', 'ay', 'al', 'h', 'z',
        'aj', 'd', 'v', 'ak', 'ba', 'n', 'j', 's', 'af', 'ax', 'at', 'aq',
        'av', 'm', 'k', 'a', 'e', 'ai', 'i', 'ag', 'b', 'am', 'aw', 'as',
        'r', 'ao', 'u', 'l', 'c', 'ad', 'au', 'bc', 'g', 'an', 'ae', 'p',
        'bb'], dtype=object)
```

```
In [40]: df_test['X0'] = le.fit_transform(df_test['X0'])
df_test['X1'] = le.fit_transform(df_test['X1'])
df_test['X2'] = le.fit_transform(df_test['X2'])
df_test['X3'] = le.fit_transform(df_test['X3'])
df_test['X4'] = le.fit_transform(df_test['X4'])
df_test['X5'] = le.fit_transform(df_test['X5'])
df_test['X6'] = le.fit_transform(df_test['X6'])
df_test['X8'] = le.fit_transform(df_test['X8'])
```

```
In [41]: df_test.head()
```

```
Out[41]:
```

	ID	X0	X1	X2	X3	X4	X5	X6	X8	X10	...	X375	X376	X377	X378	X379	X380	X382	X383	X384	...
0	1	21	23	34	5	3	26	0	22	0	...	0	0	0	1	0	0	0	0	0	0
1	2	42	3	8	0	3	9	6	24	0	...	0	0	1	0	0	0	0	0	0	0
2	3	21	23	17	5	3	0	9	9	0	...	0	0	0	1	0	0	0	0	0	0
3	4	21	13	34	5	3	31	11	13	0	...	0	0	0	1	0	0	0	0	0	0
4	5	45	20	17	2	3	30	8	12	0	...	1	0	0	0	0	0	0	0	0	0

5 rows × 377 columns

```
In [42]: df_test = df_test.drop('ID',axis=1)
```

```
In [43]: df_test.shape
```

```
Out[43]: (4209, 376)
```

```
In [44]: from sklearn.decomposition import PCA
```

```
In [45]: sklearn_pca = PCA(n_components=0.95)
```

```
In [46]: sklearn_pca.fit(df_test)
```

```
Out[46]: PCA(n_components=0.95)
```

```
In [47]: df_test=sklearn_pca.transform(df_test)
```

```
In [48]: df_test.shape
```

```
Out[48]: (4209, 6)
```

## TASK5-Predict your test\_df values using XGBoost.

```
In [49]: pip install xgboost
```

```
Requirement already satisfied: xgboost in c:\users\hp\anaconda3\lib\site-packages (1.7.1)
Requirement already satisfied: numpy in c:\users\hp\anaconda3\lib\site-packages (from xgboost) (1.21.5)
Requirement already satisfied: scipy in c:\users\hp\anaconda3\lib\site-packages (from xgboost) (1.7.3)
Note: you may need to restart the kernel to use updated packages.
```

```
In [50]: from sklearn import svm
from sklearn import model_selection
import xgboost as xgb
```

```
In [51]: model = xgb.XGBRegressor(objective="reg:linear",learning_rate=0.1)
model.fit(x_train, y_train)
y_pred = model.predict(x_test)
y_pred
model.predict(df_test)
```

```
[20:56:07] WARNING: C:/buildkite-agent/builds/buildkite-windows-cpu-autoscaling-group-i-03de431ba26204c4d-1/xgboost/xgboost-ci-windows/src/objective/regression_obj.cu:213: reg:linear is now deprecated in favor of reg:squarederror.
```



Out[51]: array([ 79.527405, 97.88469 , 98.678825, ..., 108.041794, 111.803894,  
 93.16432 ], dtype=float32)

In [ ]:

In [ ]: